

Volume II · Chapter 10 — Benchmarking & Hallucination · Labs

Source: split from `med-tech-curriculum-V1.md` (V1). From this split onward this chapter is maintained **independently** here. See `Volume-II-Chapter-10-context.md` for the AI interaction log.

Prerequisites: Ch 8, Ch 9.

Learning outcomes — the student can: evaluate LLMs on medical benchmarks (MedQA/PubMedQA/MedMCQA); interpret results; measure hallucination; build an eval harness; explain why benchmark $\neq$ clinical safety.

Labs (hands-on) — 8 h:

#	Lab	h
10a	Run a model on a MedQA-style set; compute accuracy	3
10b	Build a small eval harness with the <code>evaluate</code> library	2
10c	Hallucination test: craft adversarial clinical prompts, score	3

Datasets/tools: HF datasets/`evaluate`; local LLM. **Assessment:** eval-harness deliverable (50%); hallucination report (30%); quiz (20%). **Key decisions:** which benchmark supports which claim; honest reporting of limitations. **References:** plan §13 → *Fine-tuning, alignment & benchmarks; Medical LLMs*. **Hours:** Theory 8 + Lab 8 = 16.