

Volume IV · Chapter 20 — Structuring, Provenance & Building the Corpus · Theory

Source: split from med-tech-curriculum-V1.md (V1). From this split onward this chapter is maintained **independently** here. See Volume-IV-Chapter-20-context.md for the AI interaction log.

Prerequisites: Ch 18, Ch 19.

Learning outcomes — the student can: semantically chunk book content; attach citations/provenance; build a versioned, updatable knowledge base; manage licensing/governance of sources.

Topics (theory) — 6 h:

#	Topic	h
20.1	Semantic chunking for books/atlasses	1
20.2	Provenance model: book/edition/page/figure citations	2
20.3	Versioning & updatability of medical knowledge	1
20.4	Licensing, ownership & governance of source content	2

Datasets/tools: vector DB; chunking libraries; a metadata store. **Assessment:** cited knowledge base (**60%**); provenance audit (**20%**); quiz (**20%**). **Key decisions:** chunk granularity; provenance schema; licensing compliance; update strategy. **References:** plan §6; §13 → RAG.... **Hours:** Theory **6** + Lab **12** = **18**.